

Master of Science in Data Science

(Program Curriculum)

For Prep Sessions + Batch Start Dates:

Please refer to upgrad.com

Note: This curriculum is subject to change based on inputs from IITB and Industry



TRACK	COURSE	MODULE NAME	DESCRIPTION	SESSION
PRE-PROGRAM PREPARATORY CONTENT	Data Analysis in Excel	Data Analysis in Excel	<p><i>Taught by one of the most renowned data scientists in the country (S.Anand, CEO, Gramener), this module takes you from a beginner level Excel user to an almost professional user.</i></p>	Introduction to Excel
				Data Analysis in Excel - I: Functions, Formulae, and Charts
				Data Analysis in Excel - II: Pivots and Lookups
	Analytics Problem Solving	Analytics Problem Solving	<p><i>This module covers concepts of the CRISP-DM framework for business problem-solving.</i></p>	The CRISP-DM Framework - Business and Data Understanding
				CRISP-DM Framework - Data Preparation, Modelling, Evaluation and Deployment
	Data Analysis using SQL	Data Analysis using SQL	<p><i>Data in companies is definitely not stored in excel sheets! Learn the fundamentals of database and extract information from RDBMS using the structured query language.</i></p>	Basics of SQL: Data Retrieval, Compound Functions, Relational Operators, and Sorting
				Advanced SQL: Aggregate Functions, Nested Queries, and Joins
	Introduction to Python	Introduction to Python	<p><i>Build a foundation for the most in-demand programming language of the 21st century.</i></p>	Understanding the upGrad Coding Console
				Data Structures in Python
				Control Structure and Functions
COURSE 1: DATA TOOLKIT	Programming in Python	Programming in Python	<p><i>Learn how to approach and solve logical problems using programming.</i></p>	Logic and Syntax Building
				Data Structures: Lists, Strings, Dictionaries, and Stacks
				Time Complexity
				Searching and Sorting
				Two Pointers
	Python for Data Science	Python for Data Science	<p><i>Learn how to manipulate datasets in Python using Pandas which is the most powerful library for data preparation and analysis.</i></p>	Recursion
				Introduction to NumPy
				Operations on NumPy Arrays
				Introduction to Pandas
COURSE 2 - MACHINE LEARNING	Exploratory Data Analysis	Exploratory Data Analysis	<p><i>Learn how to find and analyse the patterns in the data to draw actionable insights.</i></p>	Getting and Cleaning Data
				Introduction to Data Visualization
				Basics of Visualization: Plots, Subplots and their Functionalities
				Plotting Data Distributions
				Plotting Categorical and Time-Series Data
	IMDb Movie Assignment	IMDb Movie Assignment	<p><i>Reinforce the concepts learnt in data science through this rigorous assignment involving the past hundred years of movie data.</i></p>	Problem Statement
				Evaluation Rubric
				Final Submission
				Solution
	Inferential Statistics	Inferential Statistics	<p><i>Build a strong statistical foundation and learn how to 'infer' insights from a huge population using a small sample.</i></p>	Data Sourcing
				Data Cleaning
				Univariate Analysis
				Segmented Univariate
				Bivariate Analysis
COURSE 3 - ADVANCED MACHİNE LEARNING	Hypothesis Testing	Hypothesis Testing	<p><i>Understand how to formulate and validate hypotheses for a population to solve real-life business problems.</i></p>	Derived Metrics
				Basics of Probability
				Discrete Probability Distributions
				Continuous Probability Distributions
				Central Limit Theorem
	Credit EDA Case Study	Credit EDA Case Study	<p><i>Solve a real industry problem through the concepts learnt in exploratory data analysis.</i></p>	Concepts of Hypothesis Testing - I: Null and Alternate Hypothesis, Making a Decision, and Critical Value Method
				Concepts of Hypothesis Testing - II: p-Value Method and Types of Errors
				Industry Demonstration of Hypothesis Testing: Two-Sample Mean and Proportion Test, A/B Testing
				Problem Statement
	Introduction to Machine Learning and Linear Regression	Introduction to Machine Learning and Linear Regression	<p><i>Venture into the machine learning community by learning how one variable can be predicted using several other variables through a housing dataset where you will predict the prices of houses based on various factors.</i></p>	Evaluation Rubric
				Final Submission
				Solution
				Simple Linear Regression
				Multiple Linear Regression
COURSE 4 - DEEP LEARNING & NLP	Linear Regression Assignment	Linear Regression Assignment	<p><i>Build a model to understand the factors car prices vary on and help a Chinese company enter the US car market.</i></p>	Industry Relevance of Linear Regression
				Problem Statement
				Evaluation Rubric
				Final Submission
	Logistic Regression	Logistic Regression	<p><i>Learn your first binary classification technique by determining which customers of a telecom operator are likely to churn versus who are not to help the business retain customers.</i></p>	Solution
				Univariate Logistic Regression
				Multivariate Logistic Regression: Model Building and Evaluation
	Unsupervised Learning: Clustering	Unsupervised Learning: Clustering	<p><i>Learn how to group elements into different clusters when you don't have any pre-defined labels to segregate them through K-means clustering, hierarchical clustering, and more.</i></p>	Logistic Regression: Industry Applications
				Introduction to Clustering
				K-Means Clustering
				Hierarchical Clustering
COURSE 5 - INDUSTRY PRACTICE	Business Problem Solving	Business Problem Solving	<p><i>Learn how to approach open ended real world problems using data as a lever to draw actionable insights.</i></p>	Other Forms of Clustering: K-Mode, K-Prototype, DB Scan
				Introduction to Business Problem Solving
				Business Problem Solving: Practical Examples
	Assignment: Unsupervised + Supervised	Assignment: Unsupervised + Supervised	<p><i>Apply the machine learning concepts learnt to solve a real-life predictive analytics problem.</i></p>	Problem Statement
				Evaluation Rubric
				Final Submission
COURSE 6 - CAREER & PLACEMENT	Case Study: Lead Scoring	Case Study: Lead Scoring	<p><i>Help the Sales team of your company identify which leads are worth pursuing through this classification case study.</i></p>	Solution
				Problem Statement
				Evaluation Rubric
COURSE 7 - CAREER & PLACEMENT	Assignment: Lead Generation	Assignment: Lead Generation	<p><i>Generate leads for your company by understanding the various lead generation techniques and applying them to a real-world scenario.</i></p>	Final Submission
				Solution
				Problem Statement

Please refer to upgrad.com

TRACK **COUR**

A horizontal bar divided into three equal-width rectangular segments. The left segment is red, the middle segment is grey, and the right segment is a lighter shade of grey.

Redacted content



Red **Grey** **White**

SPECIALISATION 1: DEEP LEARNING

COURSE 4 - DEEP LEARNING AND NEURAL NETWORKS

Model Selection & General ML Techniques	<p>Learn the pros and cons of simple and complex models and the different methods for quantifying model complexity, along with general machine learning techniques like feature engineering, model evaluation, and many more.</p>	<i>Random Forests</i>
		<i>Principles of Model Selection</i>
		<i>Model Evaluation</i>
		<i>Model Selection: Best Practices</i>
		<i>Introduction to Boosting and AdaBoost</i>
		<i>Gradient Boosting</i>
		<i>Generalized Linear Regression</i>
		<i>Regularized Regression</i>
		<i>Problem Statement</i>
		<i>Evaluation Rubric</i>
Bagging and Boosting	<p>Learn about ensemble modelling through bagging and boosting and understand how weak algorithms can be transformed into stronger ones.</p>	<i>Final Submission</i>
		<i>Solution</i>
		<i>Principal Component Analysis and Singular Value Decomposition</i>
		<i>Principal Component Analysis in Python</i>
		<i>Introduction to Time Series and its Components</i>
		<i>Working with Stationary Time Series</i>
		<i>End-to-End Analysis of Time Series</i>
		<i>Problem Statement</i>
		<i>Evaluation Rubric</i>
		<i>Final Submission</i>
Advanced Regression Assignment	<p>In this module, take a more advanced look at regression models and learn the concepts related to regularization.</p> <p>Build a regularized regression model to understand the most important variables to predict the house prices in Australia.</p>	<i>Solution</i>
		<i>Principal Component Analysis and Singular Value Decomposition</i>
		<i>Principal Component Analysis in Python</i>
		<i>Introduction to Time Series and its Components</i>
		<i>Working with Stationary Time Series</i>
		<i>End-to-End Analysis of Time Series</i>
		<i>Problem Statement</i>
		<i>Evaluation Rubric</i>
		<i>Final Submission</i>
		<i>Solution</i>
Principal Component Analysis	<p>In this module, take a more advanced look at regression models and learn the concepts related to regularization.</p>	<i>Structure of Neural Networks</i>
		<i>Feed Forward in Neural Networks</i>
		<i>Backpropagation in Neural Networks</i>
		<i>Modifications to Neural Networks</i>
		<i>Hyperparameter Tuning in Neural Networks</i>
		<i>Problem Statement</i>
		<i>Evaluation Rubric</i>
		<i>Final Submission</i>
		<i>Solution</i>
		<i>Introduction to Convolutional Neural Networks</i>
Time Series Analysis	<p>In this module, you will learn how to analyse and forecast a series that varies with time.</p>	<i>Building CNNs with Python and Keras</i>
		<i>Style Transfer and Object Detection</i>
		<i>Industry Demonstration: Using CNNs with Flowers Images</i>
		<i>Industry Demonstration: Using CNNs with X-Ray Images</i>
		<i>What Makes a Neural Network Recurrent</i>
		<i>Variants of RNNs: Bidirectional RNNs and LSTMs</i>
		<i>Building RNNs in Python</i>
		<i>Problem Statement</i>
		<i>Evaluation Rubric</i>
		<i>Final Submission</i>
Telecom Churn Case Study	<p>Solve the most crucial business problem for a leading telecom operator in India and southeast Asia - predicting customer churn.</p>	<i>Introduction to Neural Networks</i>
		<i>Neural Networks Assignment</i>
		<i>Structure of Neural Networks</i>
		<i>Feed Forward in Neural Networks</i>
		<i>Backpropagation in Neural Networks</i>
		<i>Modifications to Neural Networks</i>
		<i>Hyperparameter Tuning in Neural Networks</i>
		<i>Problem Statement</i>
		<i>Evaluation Rubric</i>
		<i>Final Submission</i>
Introduction to Neural Networks	<p>Learn the most sophisticated and cutting-edge technique in machine learning - Artificial Neural Networks or ANNs</p>	<i>Neural Networks Assignment</i>
		<i>Structure of Neural Networks</i>
		<i>Feed Forward in Neural Networks</i>
		<i>Backpropagation in Neural Networks</i>
		<i>Modifications to Neural Networks</i>
		<i>Hyperparameter Tuning in Neural Networks</i>
		<i>Problem Statement</i>
		<i>Evaluation Rubric</i>
		<i>Final Submission</i>
		<i>Solution</i>
Convolutional Neural Networks - Introduction and Industry Applications	<p>Build a neural network from scratch in Numpy to identify handwritten digits.</p>	<i>Convolutional Neural Networks - Introduction and Industry Applications</i>
		<i>Neural Networks Assignment</i>
		<i>Structure of Neural Networks</i>
		<i>Feed Forward in Neural Networks</i>
		<i>Backpropagation in Neural Networks</i>
		<i>Modifications to Neural Networks</i>
		<i>Hyperparameter Tuning in Neural Networks</i>
		<i>Problem Statement</i>
		<i>Evaluation Rubric</i>
		<i>Final Submission</i>
Recurrent Neural Networks	<p>Ever wondered what goes behind machine translation, sentiment analysis, speech recognition? Learn how RNN helps in these areas having sequential data like text, speech, videos, and a lot more.</p>	<i>Recurrent Neural Networks</i>
		<i>Neural Networks Assignment</i>
		<i>Structure of Neural Networks</i>
		<i>Feed Forward in Neural Networks</i>
		<i>Backpropagation in Neural Networks</i>
		<i>Modifications to Neural Networks</i>
		<i>Hyperparameter Tuning in Neural Networks</i>
		<i>Problem Statement</i>
		<i>Evaluation Rubric</i>
		<i>Final Submission</i>
Convolutional Neural Networks - Introduction and Industry Applications	<p>Build a neural network from scratch in Numpy to identify handwritten digits.</p>	<i>Convolutional Neural Networks - Introduction and Industry Applications</i>
		<i>Neural Networks Assignment</i>
		<i>Structure of Neural Networks</i>
		<i>Feed Forward in Neural Networks</i>
		<i>Backpropagation in Neural Networks</i>
		<i>Modifications to Neural Networks</i>
		<i>Hyperparameter Tuning in Neural Networks</i>
		<i>Problem Statement</i>
		<i>Evaluation Rubric</i>
		<i>Final Submission</i>
Recurrent Neural Networks	<p>Ever wondered what goes behind machine translation, sentiment analysis, speech recognition? Learn how RNN helps in these areas having sequential data like text, speech, videos, and a lot more.</p>	<i>Recurrent Neural Networks</i>
		<i>Neural Networks Assignment</i>
		<i>Structure of Neural Networks</i>
		<i>Feed Forward in Neural Networks</i>
		<i>Backpropagation in Neural Networks</i>
		<i>Modifications to Neural Networks</i>
		<i>Hyperparameter Tuning in Neural Networks</i>
		<i>Problem Statement</i>
		<i>Evaluation Rubric</i>
		<i>Final Submission</i>

Master of Science in Data Science

(Program Curriculum)

For Prep Sessions + Batch Start Dates:

Please refer to upgrad.com

Note: This curriculum is subject to change based on inputs from IITB and Industry



DATA SCIENCE

Master of Science in Data Science

(Program Curriculum)

For Prep Sessions + Batch Start Dates:

Please refer to upgrad.com



Note: This curriculum is subject to change based on inputs from IITB and Industry

SPECIALISATION 3: BUSINESS INTELLIGENCE/DATA ANALYTICS

COURSE 3 - SQL AND NOSQL DATABASES

TRACK	COURSE	MODULE NAME	DESCRIPTION	SESSION
		Introduction to Databases	<i>Learn how data is stored and which database is optimal to use in a particular scenario.</i>	OLAP vs OLTP Designing a database schema Creating ER Diagram
		Advanced SQL and Best Practices	<i>Apply advanced SQL concepts like windowing and procedures to derive insights from data and answer pertinent business questions.</i>	Stored functions and procedures in SQL Conditional constructs in SQL Query Optimisation
		Schema Design and Retrieval Assignment	<i>Design a schema with a predefined objective and create a database and perform analysis.</i>	Problem Statement Evaluation Rubric Final Submission Solution
		NoSQL Databases and Best Practices	<i>Take your knowledge of query languages a step further by learning about MongoDB - a NoSQL database which is becoming more and more popular in the industry.</i>	Introduction to NoSQL Database Configuring MongoDB CRUD and MongoDB Shell Indexing and Aggregation CouchDB/Redis/HBase/Neo4J
		Introduction to Cloud and Hive	<i>Understand the basics of cloud and learn about the Hive Query Language.</i>	Understanding Big Data Introduction to Apache Hive Hive Data Models - Partitions and Buckets File Formats in Apache Hive Data Analysis using Apache Hive
		SQL Case Study: Analysing Big Data in Retail	<i>Understand how a project in the industry is taken up and solved through a comprehensive business case study.</i>	Problem Statement Evaluation Rubric Final Submission Solution
		Advanced Excel	<i>Learn the advanced concepts in Excel and start to perform data analysis like a pro!</i>	Advanced Functions in Excel Slicers with Pivots Map Visualization Data Analysis Toolpak: Descriptive Summaries, Correlation, Regression
		Visualisation using Tableau	<i>Learn advanced visualisation techniques using the most in-demand visualization tool in the industry.</i>	Data Exploration in Tableau Visualising and Analysing Data in Tableau Advanced Visualisations using Tableau
		Interactive Marketing Campaign Analysis	<i>Apply the new found Excel and Tableau skills to solve an exciting business assignment.</i>	Problem Statement Evaluation Rubric Final Submission Solution
		Visualisation using PowerBI	<i>Take your visualization game a step forward by understanding how to operate PowerBI.</i>	PowerBI: Introduction and Setup Visualising and Analysing Data in PowerBI Data Transformations using PowerBI Making Interactive Dashboards with PowerBI

COURSE 4 - NATURAL LANGUAGE PROCESSING

		Introduction to R and RShiny	<i>Get a brief introduction to another popular data science language and learn how to manipulate dataframes in R and learn to create attractive dashboards and web applications using ShinyR.</i>	Introduction to R Data Structures and Programming Constructs in R Dataframe Manipulation in R Visualizations in R Creating Web Applications using RShiny
		Effective Communication Strategies, Formats, and Templates	<i>Learn how to effectively strategise, communicate, and fine grain your data analysis projects.</i>	Introduction to MS Powerpoint Identifying the Right Insights Documenting Insights Choosing the Right View for your Dashboard Storytelling with Visualization
		Presentations to Technical and Non-Technical Stakeholders	<i>Understand how to optimally present your findings to technical and non-technical stakeholders and upgrade your storytelling skills.</i>	Understanding the End Goal of your Presentation Understanding your Stakeholder and Stakeholder Empathy Levels of Details for Different Stakeholders - CXO/Leadership Vs Team Presentations Do's and Don'ts of Making a Presentation Case Study: Making a Presentation for Different Stakeholders
		Business Case Study	<i>Understand how a project in the industry is taken up and solved through a comprehensive business case study.</i>	Problem Statement Evaluation Rubric Final Submission Solution
		Capstone Project	<i>Solve an end-to-end real-life industry problem from a wide variety of domains like Marketing, Retail, E-Commerce, Supply Chain, Healthcare, BFSI, and many more.</i>	An Overview of the Domain and Associated Concepts Problem Statement Evaluation Rubric Final Submission Solution

Master of Science in Data Science

(Program Curriculum)

For Prep Sessions + Batch Start Dates:

Please refer to upgrad.com



Note: This curriculum is subject to change based on inputs from IITB and Industry

SPECIALISATION 4: BUSINESS ANALYTICS

COURSE 4 - BUSINESS REQUIREMENTS

TRACK	COURSE	MODULE NAME	DESCRIPTION	SESSION
		Tree Models	<p><i>Learn how the human decision making process can be replicated using a decision tree and other powerful ensemble algorithms.</i></p>	<i>Introduction to Decision Trees</i> <i>Algorithms for Decision Tree Construction</i> <i>Truncation and Pruning</i> <i>Random Forests</i>
		Time Series Analysis	<p><i>In this module, you will learn how to analyse and forecast a series that varies with time.</i></p>	<i>Introduction to Time Series and its Components</i> <i>Working with Stationary Time Series</i> <i>End-to-End Analysis of Time Series</i>
		Time Series Assignment	<p><i>Apply the concepts learnt in time series to solve a forecasting problem.</i></p>	<i>Problem Statement</i> <i>Evaluation Rubric</i> <i>Final Submission</i> <i>Solution</i>
		Model Selection & General ML Techniques	<p><i>Learn the pros and cons of simple and complex models and the different methods for quantifying model complexity, alongwith general machine learning techniques like feature engineering, model evaluation, and many more.</i></p>	<i>Principles of Model Selection</i> <i>Model Evaluation</i> <i>Model Selection: Best Practices</i>
		Telecom Churn Case Study	<p><i>Solve the most crucial business problem for a leading telecom operator in India and southeast Asia - predicting customer churn.</i></p>	<i>Problem Statement</i> <i>Evaluation Rubric</i> <i>Final Submission</i> <i>Solution</i>
		Advanced SQL and Best Practices	<p><i>Apply advanced SQL concepts like windowing and procedures to derive insights from data and answer pertinent business questions.</i></p>	<i>Stored functions and procedures in SQL</i> <i>Conditional constructs in SQL</i> <i>Query Optimisation</i>
		Advanced Excel	<p><i>Learn advanced visualisation techniques using the most in-demand visualization tool in the industry.</i></p>	<i>Advanced Functions in Excel</i> <i>Slicers with Pivots</i> <i>Map Visualization</i> <i>Data Analysis Toolpak: Descriptive Summaries, Correlation, Regression</i>
		Structured Problem Solving using Frameworks	<p><i>Apply the new found Excel and Tableau skills to solve an exciting business assignment.</i></p>	<i>Introduction to Structured Problem Solving</i> <i>Understanding the Business Problem - I: Interviewing and Basic Frameworks</i> <i>Understanding the Business Problem - II: 5W, 5WHYS, SPIN Framework, etc.</i>
		Hypothesis Formulation	<p><i>The module will equip you with a stepwise process for understanding a business problem and building hypotheses around it.</i></p>	<i>Introduction to Hypothesis Formulation</i> <i>Issue Tree Framework</i> <i>Specialized Frameworks: 4Ps and 5C</i>
		Assignment	<p><i>Apply your learnings from the course to solve a real-life business problem.</i></p>	<i>Problem Statement</i> <i>Evaluation Rubric</i> <i>Final Submission</i> <i>Solution</i>
		Revenue and Operational Cost Modelling	<p><i>Learn about evaluation models covering revenue and operational cost modelling.</i></p>	<i>Introduction to Cost Modelling</i> <i>Modeling Frameworks</i> <i>Revenue and Operational Cost Modelling in Various Domains</i>
		Introduction to Economics and Financial Concepts	<p><i>Understand financial concepts such as revenue, cost of goods sold, profit, balance sheets, cash flow statements, and EBITDA. Also, get a brief introduction to economics concepts including supply and demand curves, cost curves, and a lot more!</i></p>	<i>Introduction to Financial Concepts</i> <i>Revenue, Balance Sheets, Cash Flow Statements</i> <i>EBITDA</i> <i>Introduction to Economics</i> <i>Supply-Demand Curves</i> <i>Basic Econometrics</i>
		Effective Communication Strategies, Formats, and Templates	<p><i>Learn how to effectively strategise, communicate, and fine grain your data analysis projects.</i></p>	<i>Introduction to MS Powerpoint</i> <i>Identifying the Right Insights</i> <i>Documenting Insights</i> <i>Choosing the Right View for your Dashboard</i> <i>Storytelling with Visualization</i>
		Presentations to Technical and Non-Technical Stakeholders	<p><i>Understand how to optimally present your findings to technical and non-technical stakeholders and upgrade your storytelling skills.</i></p>	<i>Understanding the End Goal of your Presentation</i> <i>Understanding your Stakeholder and Stakeholder Empathy</i> <i>Levels of Details for Different Stakeholders - CXO/Leadership Vs Team Presentations</i> <i>Do's and Don'ts of Making a Presentation</i> <i>Case Study: Making a Presentation for Different Stakeholders</i>
		Business Case Study	<p><i>Understand how a project in the industry is taken up and solved through a comprehensive business case study.</i></p>	<i>Problem Statement</i> <i>Evaluation Rubric</i> <i>Final Submission</i> <i>Solution</i>
		Capstone Project	<p><i>Solve an end-to-end real-life industry problem from a wide variety of domains like Marketing, Retail, E-Commerce, Supply Chain, Healthcare, BFSI, and many more.</i></p>	<i>An Overview of the Domain and Associated Concepts</i> <i>Problem Statement</i> <i>Evaluation Rubric</i> <i>Final Submission</i> <i>Solution</i>

Master of Science in Data Science

(Program Curriculum)

For Prep Sessions + Batch Start Dates:

Please refer to upgrad.com

Note: This curriculum is subject to change based on inputs from IITB and Industry



SPECIALISATION 5: DATA ENGINEERING

COURSE 3 - DATA ENGINEERING - I

TRACK	COURSE	MODULE NAME	DESCRIPTION	SESSION
		Introduction to Hadoop and MapReduce Programming	<i>Understand the world of distributed data processing and storage with Hadoop. Learn to write MapReduce jobs in Python.</i>	Concepts related to distributed computing Hadoop Distributed File System MapReduce Programming in Python
		Data Management and Relational Database Modelling	<i>Understand the concepts of Data Management and learn to model data for a Relational Database.</i>	Enterprise Data Management Relational Database Modelling Normal Forms and ER Diagrams
		NoSQL Databases and Apache HBase	<i>Learn the concepts of NoSQL databases. Understand the workings of Apache HBase.</i>	Concepts of NoSQL Databases Introduction to Apache HBase HBase Python API Comparision of NoSQL Databases
		Data Warehousing (Optional)	<i>Understand the intricacies behind designing a data warehouse and a data lake for your use case.</i>	Introduction to Data Warehouse and Data Lakes Designing Data Warehousing for an ETL Data Pipeline Designing Data Lake for an ETL Data Pipeline
		Data Ingestion with Apache Sqoop and Apache Flume	<i>Get familiar with the challenges involved in data ingestion. Use Sqoop and Flume to ingest structured and unstructured data into Hadoop.</i>	Introduction to Data Ingestion Structured data ingestion with Sqoop Unstructured data ingestion with Flume
		Assignment (Optional)	<i>Practice the concepts learnt so far with this comprehensive assignment.</i>	Problem Statement Evaluation Rubric Final Submission Solution
		Building and Querying Data Warehouse with Apache Hive	<i>Manage and query a data warehouse with Apache Hive. Learn to write optimized HQL for large scale data analysis.</i>	Fundamentals of Apache Hive Writing HQL for Data Analysis Hive Query Optimization
		Case Study: Ingestion & Warehousing	<i>Make use of Sqoop, Flume, Hive and HBase to design an ETL data pipeline.</i>	Introduction and Problem Statement Grading Rubrics and Submission
		Data Processing with PySpark	<i>Get introduced to Apache Spark, a lightning fast big data processing engine. Use PySpark to create large scale data processing applications.</i>	Introduction to Apache Spark Apache Spark Architecture PySpark APIs - RDDs, DataFrames, SQL Spark Job optimization
		Real-Time Data Streaming with Apache Kafka	<i>Understand the producer-consumer architecture of Apache Kafka. Learn to set up a Kafka cluster for managing real-time data.</i>	Fundamentals of Apache Kafka Setting up Kafka Producer and Consumer Kafka Connect API
		"Real-Time Data Processing using Spark Streaming"	<i>Learn about the real-time data processing architecture of Apache Spark. Build Spark Streaming applications to process data in real-time.</i>	Spark Streaming Architecture Spark Streaming APIs Building Stream Processing Application with Spark
		Assignment (Optional)	<i>Use Kafka and Spark to develop a real-time data processing application.</i>	Problem Statement Evaluation Rubric Final Submission Solution
		Building Automated Data Pipelines with Oozie/Airflow	<i>Automate your Data Pipeline with Apache Airflow.</i>	Fundamentals of Oozie/Airflow Workflow Management with Oozie/Airflow Automating an entire Data Pipeline with Oozie/Airflow
		Analytics using PySpark	<i>Use PySpark to do EDA and Predictive Analysis of large datasets.</i>	"Exploratory Data Analysis with PySpark" Predictive Analysis with Spark MLlib
		Case Study: Kafka, Spark Streaming and PySpark	<i>Build an end-to-end real-time data processing application using Spark Streaming and Kafka.</i>	Introduction and Problem Statement Grading Rubrics and Submission

COURSE 4 - DATA ENGINEERING - II

	Capstone Project	<i>The capstone project will stitch all the components of data engineering together.</i>	An Overview of the Domain and Associated Concepts Problem Statement Evaluation Rubric Final Submission Solution
--	-------------------------	--	---

Master of Science in Data Science

(Program Curriculum)

For Prep Sessions + Batch Start Dates:

Please refer to upgrad.com



Note: This curriculum is subject to change based on inputs from IIITB and Industry

RESEARCH METHODOLOGY

(DURATION: 2 MONTHS)

10 CREDITS

WHAT IS RESEARCH?

Familiarise with different aspects of research

- Intro to research
- Importance of research
- Criticism in research and its importance
- Peer reviews in research and its importance

TYPES OF RESEARCH

Develop an understanding of various research design and techniques

- Descriptive vs Analytical
- Applied vs Fundamental
- Quantitative vs Qualitative
- Bayesian vs Frequentist Approach

RESEARCH PROCESS

Learn about the different steps in the research process and how to evaluate a literature

- Research question
- Hypothesis and aims
- Formulating a Problem
- Literature review

RESEARCH PROJECT MANAGEMENT

Learn how to plan the project timelines and arrange for data & software

- Understand the different steps involved in a project cycle
- Project Requirements on Data
- Identifying the milestones in a research project
- Learn how to track the progress using Gantt Chart

REPORT WRITING AND PRESENTATION

Master good scientific writing and proper presentation skills

- Art of writing papers
- Parts of a paper
- Tools to write papers
- Publishing papers: Journals + Seminars

SCIENTIFIC ETHICS

Develop an understanding of the ethical dimension in research

- Citation Methods and Rules
- Honor Code
- Research Claims
- IPh

MASTER'S DISSERTATION

(DURATION: 4 MONTHS)

60 CREDITS

MASTER'S THESIS

Articulate, research and present your project.

- Monthly Checkpoints
- Submission

An example of project outlines is here:

- Investigate the risk factors for eye disease from complex longitudinal datasets
- Investigate a diagnosis of eye diseases using imaging ophthalmic data
- Multi-task learning for drug design and discovery
- Using stacking for brain tumour discrimination
- Investigate dietary patterns and metabolite fingerprints of takeaway (fast) food consumers using PCA and Clustering methods
- Longitudinal studies to investigate the complex link between corporate environment engagement, green disclosure, business model transformation and supply chain performance
- Preventing credit card fraud through pattern recognition
- Developing a recommender system for a Media giant
- Using social media feed to place tweets regarding natural disasters on a map